

CS47100

Introduction to Artificial Intelligence

Lecture 39: Final Review

Dec 6th, 2024

Jiaxin and Medha (Purdue)

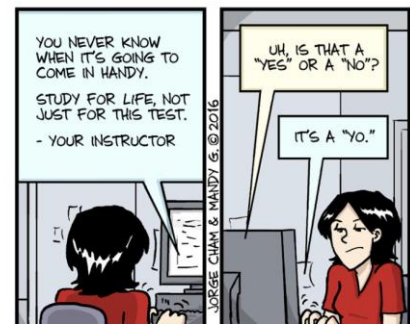
CS47100 (Fall 2024)

1

1

Final Exam

- **Time:**
 - Thu Dec 12, 8:00 am – 10:00 am
- **Location:**
 - BHEE 129
- **What to bring:**
 - Cheat Sheet
 - Calculator
 - University ID
 - Pen, Pencil



WWW.PHDCOMICS.COM

Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

2

Final Exam

- **6 Problems (Subject to change, but small variations):**

- First order logic
- Bayes Net
- MDP
- RL
- Supervised Learning
- Multiple Choice

Exam Tips

- Answer box's size does not correlate with the solution
- Don't have to write the exam in order
 - Its order is based on course content
- Start with the content you are most familiar with
- Study the examples covered in lecture & homework

What to write on the page of note?

Recommend:

- Definitions
- Equations
- Algorithms / Procedure to a type of problem
- General facts/summary about a topic
- Organize the note based on topics

Not recommended:

- Print this slide into one page...
- Putting homework problems/solutions
- Writing too much (Too difficult to find during an exam)



First order logic

Basic syntactic elements of FOL

• Constants	John, Richard, 2,...	Constants: objects
• Predicates	Brother, >,...	Relation: set of tuples of objects that are related e.g., {<Richard, John>, <John, Richard>}
• Functions	Sqrt, LeftLegOf,...	
• Variables	x, y, a, b,...	Function: relations where a given set of objects are related to exactly one object e.g., {<John>→John's left leg, <Richard>→Richard's left leg}
• Connectives	$\neg, \Rightarrow, \wedge, \vee, \Leftrightarrow$	
• Equality	=	
• Quantifiers	\forall, \exists	

Universal quantification

- $\forall \langle \text{variables} \rangle \langle \text{sentence} \rangle$
- Everyone in CS471 is smart: $\forall x \text{ in}(x, \text{CS471}) \Rightarrow \text{Smart}(x)$
- $\forall x$ P is true in a model m iff P is true with x interpreted as *each* possible object in the model
- Roughly speaking, equivalent to the **conjunction** of **instantiations** of P

$\text{In}(\text{John}, \text{CS471}) \Rightarrow \text{Smart}(\text{John})$
 $\wedge \text{In}(\text{Jane}, \text{CS471}) \Rightarrow \text{Smart}(\text{Jane})$
 $\wedge \text{In}(\text{CS471}, \text{CS471}) \Rightarrow \text{Smart}(\text{CS471})$
 $\wedge \dots$

Existential quantification

- $\exists \langle \text{variables} \rangle \langle \text{sentence} \rangle$
- Someone in CS471 is smart: $\exists x \text{In}(x, \text{CS471}) \wedge \text{Smart}(x)$
- $\exists x P$ is true in a model m iff P is true with x interpreted as *some* possible object in the model
- Roughly speaking, equivalent to the **disjunction** of **instantiations** of P

$\text{In}(\text{John}, \text{CS471}) \wedge \text{Smart}(\text{John})$
 $\vee \text{In}(\text{Jane}, \text{CS471}) \wedge \text{Smart}(\text{Jane})$
 $\vee \text{In}(\text{CS471}, \text{CS471}) \wedge \text{Smart}(\text{CS471})$
 $\vee \dots$

Unification

- Two sentences α, β can be **unified** with substitution θ if
 - $\text{SUBST}(\theta, \alpha) = \text{SUBST}(\theta, \beta)$
 - $\theta = \{x/v\}$ is a substitution of variable x with v

α	β	θ
$\text{Knows}(\text{John}, x)$	$\text{Knows}(\text{John}, \text{Jane})$	$\{x/\text{Jane}\}$
$\text{Knows}(\text{John}, x)$	$\text{Knows}(y, \text{Steve})$	$\{x/\text{Steve}, y/\text{John}\}$
$\text{Knows}(\text{John}, x)$	$\text{Knows}(y, \text{Mother}(y))$	$\{y/\text{John}, x/\text{Mother}(\text{John})\}$
$\text{Knows}(\text{John}, x)$	$\text{Knows}(x, \text{Steve})$	$\{\text{fail}\}$

Resolution

- **First convert to CNF** (as with propositional logic)
 - Conjunction clauses of disjunction terms
 - $(\text{Term}_1 \vee \text{Term}_2 \vee \text{Term}_3) \wedge (\text{Term}_4 \vee \text{Term}_5)$
- **Negate the conclusion α** and add it to the KB, i.e., add $\neg\alpha$ to KB
- **Resolution Rule**

$$\frac{\ell_1 \vee \dots \vee \ell_k, \quad m_1 \vee \dots \vee m_n}{\ell_1 \vee \dots \vee \ell_{i-1} \vee \ell_{i+1} \vee \dots \vee \ell_k \vee m_1 \vee \dots \vee m_{j-1} \vee m_{j+1} \vee \dots \vee m_n}$$
- where ℓ_i and m_j are **complementary** given a **substitution**.
 - i.e., $\neg\text{SUBST}(\theta, \ell_i) = \text{SUBST}(\theta, m_j)$
- Apply resolution and **find contradiction**, i.e., arrive at an empty clause

Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

11

11

Resolution --- Example

- $\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$
- $\text{King}(\text{John})$
- $\forall x \text{ Greedy}(x) = \forall y \text{ Greedy}(y)$ (**Standardizing apart**: eliminates overlap of variables)
- $\text{Brother}(\text{Richard}, \text{John})$
- $\alpha: \text{Evil}(\text{John})$

$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha)$	commutativity of \wedge
$(\alpha \vee \beta) \equiv (\beta \vee \alpha)$	commutativity of \vee
$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma))$	associativity of \wedge
$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma))$	associativity of \vee
$\neg(\neg\alpha) \equiv \alpha$	double-negation elimination
$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha)$	contraposition
$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta)$	implication elimination
$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha))$	biconditional elimination
$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta)$	De Morgan
$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta)$	De Morgan
$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma))$	distributivity of \wedge over \vee
$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma))$	distributivity of \vee over \wedge

Convert to CNF

- $\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$
 - $\neg(\text{King}(x) \wedge \text{Greedy}(x)) \vee \text{Evil}(x)$ (**Implication Elimination**)
 - $\neg\text{King}(x) \vee \neg\text{Greedy}(x) \vee \text{Evil}(x)$ (**De Morgan**)

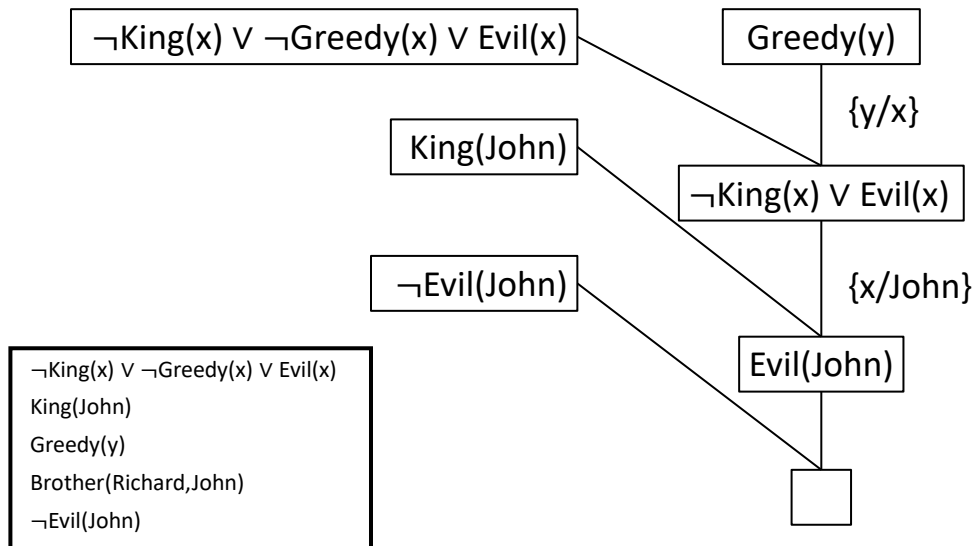
Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

12

12

Resolution --- Example



Probability

Probabilities

Probability space (Ω, \mathcal{F}, P)


- **Sample Space** Ω (Possible outcomes)
- \mathcal{F} = **Events** = a set of subsets of Ω
 - An **event** is a subset of Ω
- P = **Probability measure** on \mathcal{F}
- $P(A)$ = Probability of event A .

Dice roll:

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$


$$\text{Event } A = \text{Getting } 1 = \{1\}. \quad P(A) = P(\{1\}) = \frac{1}{6}$$

$$\text{Event } B = \text{Getting } 1 \text{ or } 3 = \{1, 3\} \quad P(B) = P(\{1, 3\}) = \frac{2}{6}$$




Expectations

- The **expected value** of a random variable is the average, weighted by the probability distribution over outcomes
- **Example:** How long to get to the airport?



Independence

- Two variables are **independent** if:
 - This says that their joint distribution factors into a product two simpler distributions $\forall x, y : P(x, y) = P(x)P(y)$
 - Another form: $\forall x, y : P(x|y) = P(x)$
 - We write: $X \perp\!\!\!\perp Y$
- Independence is a simplifying modeling assumption
 - Empirical joint distributions at best “close” to independent
 - What could we assume for {Weather, Traffic, Cavity, Toothache}?



Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)


17

17

The Product Rule

- Sometimes have conditional distributions but want the joint distribution

$$P(y)P(x|y) = P(x, y) \quad \longleftrightarrow \quad P(x|y) = \frac{P(x, y)}{P(y)}$$



Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

18

18

The Chain Rule

- More generally, can always write any joint distribution as an incremental product of conditional distributions

$$P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_1, x_2)$$

$$P(x_1, x_2, \dots, x_n) = \prod_i P(x_i|x_1 \dots x_{i-1})$$

- Why is this always true?
 - Repeated application of product rule

Bayes' Rule

- Two ways to factor a joint distribution over two variables:

$$P(x, y) = P(x|y)P(y) = P(y|x)P(x)$$

- Dividing $P(y)$, we get:

$$P(x|y) = \frac{P(y|x)}{P(y)}P(x)$$

- Why is this at all helpful?
 - Let's us build one conditional from its reverse
 - Often one conditional is tricky but the other one is simple
 - Foundation of many systems we'll see later



Reverend Thomas Bayes

Bayesian Network

Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

21


21

Bayesian Network (BN) Semantics

- A set of nodes, one per random variable X
- A directed, acyclic graph
- A conditional distribution for each node
 - A collection of probability distributions over X , one **for each combination** of parents' values

$$P(X|a_1 \dots a_n)$$

- CPT: conditional probability table
- Description of a noisy “causal” process



$$P(X|A_1 \dots A_n)$$

Bayesian network =
Topology (graph) + Local
Conditional Probabilities

Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

22

22

Probabilities in BNs


- Why are we guaranteed that BN results in a proper joint distribution?

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i))$$


- Chain rule (valid for all distributions): $P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_1 \dots x_{i-1})$
- Assume conditional independence: $P(x_i | x_1, \dots, x_{i-1}) = P(x_i | \text{parents}(X_i))$
 - Consequence: $P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i))$
- A BN **cannot** represent all possible joint distributions
 - The topology **enforces** certain conditional independencies

Active / Inactive Paths

- Question:** Are X and Y conditionally independent given evidence variables {Z}?
 - Yes, if X and Y “**d-separated**” by Z
 - Consider all (undirected) paths from X to Y
 - No active paths = independence!
- A path is **active** if each triple is active:
 - Causal chain $A \rightarrow B \rightarrow C$ where **B is unobserved** (either direction)
 - Common cause $A \leftarrow B \rightarrow C$ where **B is unobserved**
 - Common effect (aka v-structure) $A \rightarrow B \leftarrow C$ where **B or one of its descendants is observed**
- In a path, if one triple is **inactive**, the entire path is inactive



D-Separation



- Query: $X_i \perp\!\!\!\perp X_j | \{X_{k_1}, \dots, X_{k_n}\} ?$
- Check all (undirected!) paths between X_i and X_j
 - If one or more active, then independence **not guaranteed**
- Otherwise (i.e., if all paths are **inactive**), then **independence is guaranteed**

$$X_i \not\perp\!\!\!\perp X_j | \{X_{k_1}, \dots, X_{k_n}\} \text{ Maybe?}$$

$$X_i \perp\!\!\!\perp X_j | \{X_{k_1}, \dots, X_{k_n}\}$$

Simple sampling

- Given an empty network (no evidence)
 - Begin with nodes without parents
 - Sample from CPDs sequentially to instantiate all nodes
- This will produce one sample from the joint distribution
- Do this many times to produce an empirical distribution that approximates the full joint distribution.

Rejection sampling

- Sample the network as before
 - but **discard samples** that don't correspond with the evidence.
- Similar to real-world estimation procedures, which use observation
- However, it is hopelessly expensive for large networks where $P(e)$ is small
 - **Very few observations consistent with evidence...**


Likelihood weighting

- Do simple sampling as before...
 - But only generate samples that are consistent with the evidence
 - And **weight the likelihood** of each sample based on the evidence
- More efficient than rejection-based sampling since we use all the samples generated
- ... but performance degrades as number of evidence variables increases
- Cannot deal with complex evidence (e.g., rain or sprinkler on).

Markov Decision Process

Markov Decision Processes

- An **MDP** is defined by:
 - Set of states $s \in S$
 - Set of actions $a \in A$
 - Transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - Reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - Start state s_0
 - Maybe a terminal state
- MDPs are **non-deterministic** search problems



Values of States


- Fundamental operation: compute the (expectimax) value of a state
 - Expected utility under optimal action
 - Average sum of (discounted) rewards
 - This is just what expectimax computed

- Recursive definition of value (Bellman Equation):

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^*(s')]$$



Computing Actions from Values

- Let's imagine we have the optimal values $V^*(s)$
- How should we act?
 - Not so obvious...

- We need to do a mini-expectimax (one step)

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^*(s')]$$

0.95 ▶	0.96 ▶	0.98 ▶	1.00
▲ 0.94		◀ 0.89	-1.00
▲ 0.92	◀ 0.91	◀ 0.90	◀ 0.80

- This is called **policy extraction**, since it gets the policy implied by the values

Policy Iteration

- **Evaluation:** For fixed current policy π , find values with policy evaluation:

- Solve n linear equations with n unknowns
- Iterate until values converge:

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} P(s'|s, \pi_i(s)) [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

- **Improvement:** For fixed values, get a better policy using policy extraction

- One-step look-ahead:

$$\pi_{i+1}(s) = \operatorname{argmax}_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

Reinforcement Learning

Reinforcement Learning

- **Assume** a Markov decision process (MDP):

- A set of states S
- A set of actions (per state) A
- A model $T(s, a, s')$
- A reward function $R(s, a, s')$



- Still looking for a policy $\pi(s)$


- **New twist: T or R is unknown**

- I.e., we don't know which states are good or what the actions do
- Must **try** out actions and states **to learn**

Passive Reinforcement Learning

- **Simplified task:** policy evaluation

- **Input:** a **fixed** policy $\pi(s)$
- You don't know the transitions $T(s, a, s')$
- You don't know the rewards $R(s, a, s')$
- **Goal:** learn the state values $V^\pi(s)$




- **In this case:**

- Learner is "along for the ride"
- **No choice about what actions to take**
- Just execute the policy and learn from experience
- This is **NOT** offline planning. Agent **must take** actions in the world.

Temporal Difference (TD) Learning

- **Big idea:** learn from every experience
 - Update $V(s)$ each time we experience a transition (s, a, s', r)
 - Likely outcomes s' will contribute updates more often
- **Temporal difference learning** of values
 - Policy still fixed, still doing evaluation!
 - Move values toward value of whatever successor occurs: **running average**



Sample of $V(s)$: $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to $V(s)$: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

Same update: $V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$ "Temporal Difference"

Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

37

37

Q-Learning

- Q-Learning: sample-based Q-value iteration

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$$

- Learn $Q(s, a)$ values as you go

- Receive a sample (s, a, s', r)
- Consider your old estimate: $Q(s, a)$
- Consider your new sample estimate:

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

- Incorporate the new estimate into a running average:

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha)[sample]$$

this is TD Q-learning



Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

38

38

Active Reinforcement Learning

- **Full Reinforcement learning:** find optimal policies (like value iteration)

- You don't know the transitions $T(s, a, s')$
- You don't know the rewards $R(s, a, s')$
- You choose the actions now
- **Goal: learn the optimal policy / values**



- **In this case:**

- Learner **makes** choices
- Fundamental tradeoff: **exploration vs. exploitation**
- **This is NOT offline planning.** You must take actions in the world and find out what happens...

How to Explore?

- Several schemes for forcing exploration
 - Simplest: random actions (**ϵ -greedy**)
 - Every time step, flip a coin
 - With (small) probability ϵ , act randomly
 - With (large) probability $1-\epsilon$, act on current policy
- **Problems with random actions?**
 - You do eventually explore the space but **keep acting randomly** once learning is done..
 - **One solution:** lower ϵ over time
 - Another solution: **exploration functions**



Exploration Functions



- **When to explore?**
 - **Random** actions: explore a fixed amount
 - **Better idea**: explore areas whose value has **not (yet) established**, eventually stop exploring
- **Exploration function**
 - Takes a value estimate u and a visit count n , and returns an **optimistic utility** ($k > 0$), e.g., $f(u, n) = u + k/n$


Regular Q-Update: $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha[R(s, a, s') + \gamma \max_{a'} Q(s', a')]$

Modified Q-Update: $Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha [R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))]$

- **Note:** this propagates the “**bonus**” back to states that lead to unknown states as well!

Generalizing Across States

- Basic Q-Learning keeps a table of all Q-values
- In realistic situations, **we cannot possibly learn about every single state!**
 - Too many states to visit them all in training
 - Too many states to hold the Q-tables in memory
- Instead, we want to **generalize**:
 - Learn about some small number of training states from experience
 - Generalize that experience to **new, similar situations**
 - This is the fundamental idea in machine learning



Linear Value Functions

- Using a feature representation, we can write a q function (or value function) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- Advantage:** our experience is summed up in a few powerful numbers

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Disadvantage:** states may share features but are very different in value!

Approximate Q-Learning

- Q-learning with linear Q-functions:

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

transition = (s, a, r, s')


$$\text{difference} = \left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$

Exact Q's

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$$

Approximate Q's



- Intuitive interpretation:**

- Adjust weights of active features, e.g., if something unexpectedly bad happens, blame the features that were on: “dis-prefer” all states with that state’s features

Supervised Learning

K-Nearest Neighbor Algorithm

- **Classification Model:**

- **Training set** $\mathcal{D} = \{(x^{(n)}, y^{(n)})\}_{n=1}^N$

- In Nearest Neighbor, $\theta = \mathcal{D}$ (memorize the training set!)

- Given a “distance metric” $d(x, x')$

- Find the K closest example to x , denoted as $\{n\} = S_K(x, \mathcal{D})$

- **Model:**

- $f_{\theta}(y|x) = \arg \max_y P(Y = y|x)$

- $P(Y = y|x; D) = \frac{1}{K} \sum_{n \in S_K} \mathbf{1}[y^{(n)} = y]$

$\mathbf{1}(\text{statement}) = 1$ if
statement is true, else 0

Naive Bayes --- Prediction


- **Naïve Bayes**: features are **conditionally independent** given target

- How to make a prediction?

- Compute the probability $\operatorname{argmax}_y P(Y = y | x_1, x_2, \dots)$

- How to compute $P(Y | x_1, x_2, \dots)$?

- $P(Y | x_1, x_2, \dots) = P(Y, x_1, x_2, \dots) / P(x_1, x_2, \dots)$
 $\propto P(Y, x_1, x_2, \dots)$




- $P(Y, x_1, x_2, \dots) = P(Y) \prod_k P(x_k | Y)$

Laplace's estimate (extended)

- Pretend you saw every outcome k extra times

$$P_{LAP,k}(x) = \frac{c(x) + k}{N + k|X|}$$



- What's Laplace with $k = 0$? $P_{LAP,0}(X) =$

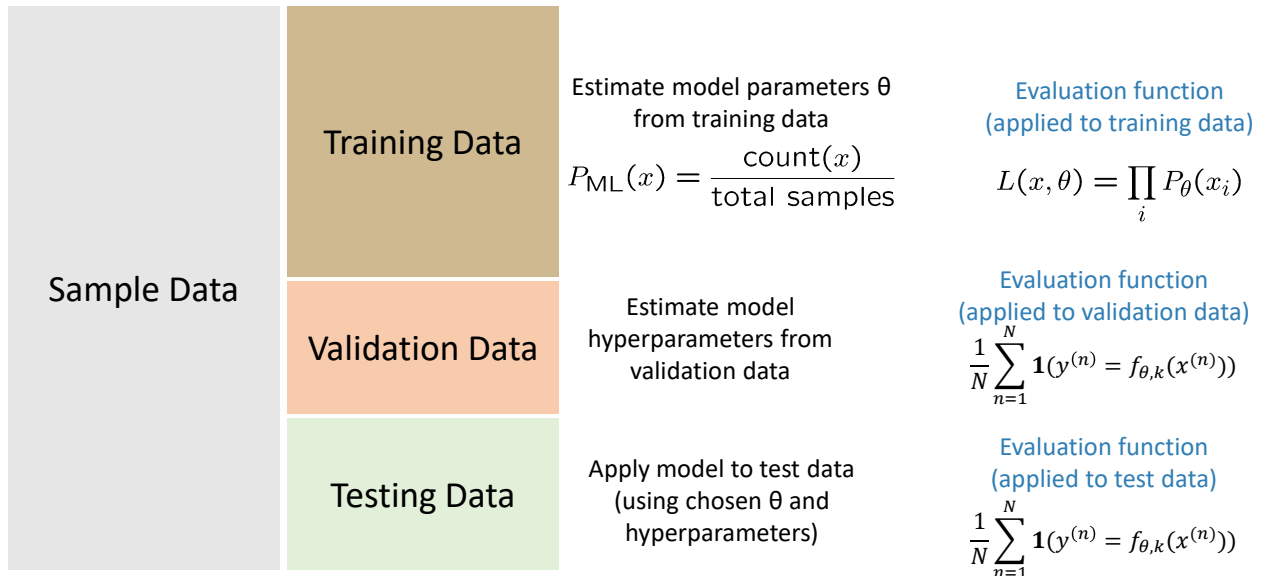
- k is the strength of the prior $P_{LAP,1}(X) =$ $P_{LAP,100}(X) =$

- Laplace for conditionals:

- Smooth each condition separately:

$$P_{LAP,k}(x|y) = \frac{c(x, y) + k}{c(y) + k|X|}$$

Training/validation/test cycle




Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

49

49

S-fold cross validation



White: training data; red:
validation data.

Every data point serves in the training and the validation dataset:

- Split data into **S** equal parts.
- Use each part **in turn** as a validation dataset and use the others as a training dataset.
- Choose the hyperparameter leading to **best average performance**.
- **Leave-one-out cross validation**: every data point is used as the validation once.
 - I.e., $S = |\mathcal{D}|$

Jiaxin and Medha (Purdue)

CS47100 (Fall 2024)

50

50

Last slide!

- **Course Evaluation**
 - Survey End Date: 12/8/2024
- **Time:**
 - Thu 12/12 08:00AM - 10:00AM
- **Location:**
 - BHEE 129
- **What to bring:**
 - Cheat Sheet
 - Calculator
 - University ID
 - Pen, Pencil

